

Artificial Intelligence Principles

6G7V0011 - 1CWK100

Dr. Peng Wang

Email: p.wang@mmu.ac.uk

Department of Computing and Mathematics

Tuesday, Nov. 12th, 2024


Outline

Search Algorithms
 Quiz and Recap
 Heuristics
Search Algorithms
 Informed Search

Outline

- Search Algorithms
 - Quiz and Recap
 - Heuristics
- Search Algorithms
 - Informed Search

Figure 3.10



Part of the Romania state space, selected to illustrate uniform-cost search.

Figure 1: Map for UCS Quiz

- Q1: Could you draw the search tree using UCS for this partial Romanian map?
- Q2: Could you write down how openset and closedset being changed?

UCS Quiz and Recap - continued

O = {'Sibius':0},
C = {}.

O = {},
C = {'Sibius':0},
current = 'Sibius'.

O = {'Fagaras':99, 'Rimnicu Vilcea':80},
C = {'Sibius':0}.

O = {'Fagaras':99},
C = {'Sibius':0, 'Rimnicu Vilcea':80}, current = 'Rimnicu Vilcea'.

O = {'Fagaras':99, 'Pitesti':177},
C = {'Sibius':0, 'Rimnicu Vilcea':80}.

O = {'Pitesti':177},
C = {'Sibius':0, 'Rimnicu Vilcea':80, 'Fagaras':99},
current='Fagaras'.

UCS Quiz and Recap - continued

O = { 'Pitesti':177, 'Bucharest':310 },
C = { 'Sibius':0, 'Rimnicu Vilcea':80, 'Fagaras':99 }.

O = { 'Bucharest':310 },
C = { 'Sibius':0, 'Rimnicu Vilcea':80, 'Fagaras':99, 'Pitesti':177 },
current='Pitesti'.

O = { 'Bucharest':278 },
C = { 'Sibius':0, 'Rimnicu Vilcea':80, 'Fagaras':99, 'Pitesti':177 }.

O = { },
C = { 'Sibius':0, 'Rimnicu Vilcea':80, 'Fagaras':99, 'Pitesti':177, 'Bucharest':278 },
current=Bucharest.




Figure 2: Same cost for each edge




Figure 3: Different cost for each edge

Recap: In uninformed search, including BFS, DFS, and UCS

- Agents do not have goal information, i.e., no idea where the *goal* is
- Agents try to minimise the cumulative path cost from *start* to *goal*
- Given a *node* n , agents evaluate the path cost by the evaluation function $f(n)$, which can be distance, time, etc. that you 'care' the most!

- **Question:** $f(n)$ for BFS (depth of node), DFS (negative/inverse/reciprocal of the depth), and UCS (path cost)?

We now define the (cumulative) path cost of a *node* n as $g(n)$, then for uninformed search, we have

$$f(n) = g(n) \tag{1}$$

What if

- Agents still do not know any goal information
- But, as designers, we can provide our experience (expert knowledge) of the environment to the agents?

Example:




Figure 4: Different heuristics for a grid environment

Given two points \mathbf{s} and \mathbf{g} of dimension N , we have

Euclidean distance

$$E(\mathbf{s}, \mathbf{g}) = \sqrt{\sum_{i=0}^N (s_i - g_i)^2}$$

Manhattan distance

$$M(\mathbf{s}, \mathbf{g}) = \sum_{i=0}^N |s_i - g_i|$$

1. Heuristic 1: $h_e(start)$ ($h_1(\mathbf{s})$): The Euclidean distance between *start* and *goal* is $5 * dx$, or simply 5 in cases where $f(n)$ is calculated same way, though it crosses the obstacles
2. Heuristic 2: $h_m(start)$ ($h_2(\mathbf{s})$): The Manhattan distance between *start* and *goal* is $7 * dx$, or simply 7. Manhattan distance is tricky, 7 corresponds to more than one path.

When there are more choices, heuristic becomes informative!




Figure 5: Heuristics when there are more than one choices

Heuristic function $h(n)$

Functions such as $h_e(n)$ or $h_m(n)$ that helps estimate the cost of the cheapest path from the state at node n to a goal state.

What we can observe from the previous Figure:

1. $n_1 \leftarrow \text{node}_1, n_2 \leftarrow \text{node}_2$
2. $g(n_1) = 1, g(n_2) = 1$
3. $f(n_1) = g(n_1) = 1, f(n_2) = g(n_2) = 1$
4. $h_e(n_1) = \sqrt{18} < h_e(n_2) = \sqrt{20}$

Since $h_e(n_1) < h_e(n_2)$, can we simply choose n_1 over n_2 because it is 'closer' to the goal according to our experiences (heuristic)?

Yes, we can! Greedy best-first search.

What is it?

- A form of best-first search that expands first the node with the lowest $h(n)$ value
- Node n appears to be closest to the goal (?)
- Likely leads to a solution quickly (?)
- Evaluation function $f(n) = h(n)$
 - No path cost counted

Greedy Best-first Search

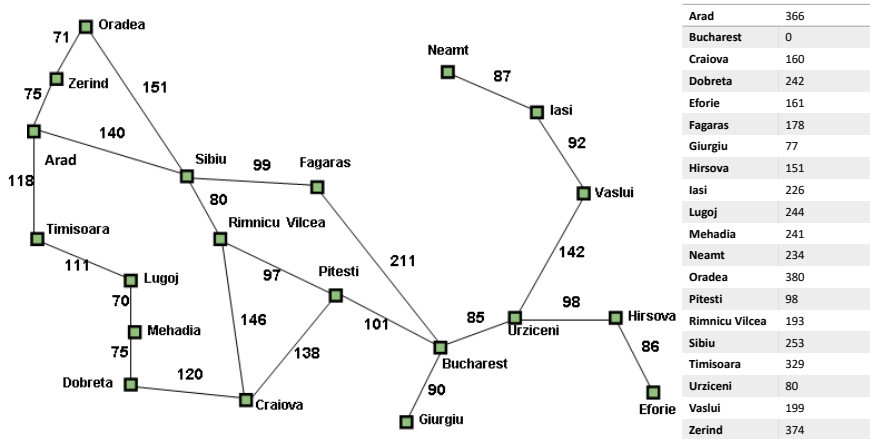


Figure 6: Heuristics on the Romania routing problem

Heuristic $h(n)$: Straight-line distance

Greedy Best-first Search

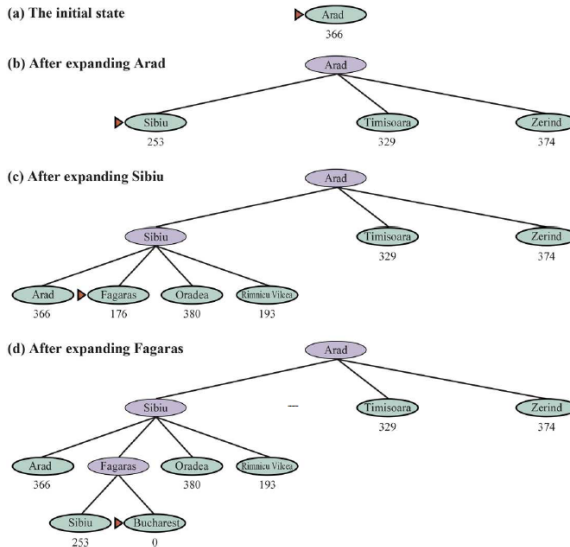




Figure 7: How greedy best-first search works on Romanian problem

Greedy Best-first Search

Strategy: always expand a node that 'you' tell the agent is closest to a goal state. Not necessarily the best solution (Romania problem).



- Takes you straight to the goal, with a sub-optimal solution



- Like a badly-guided DFS, perfectly misses goal(s)

Outline

Search Algorithms
Quiz and Recap
Heuristics
Search Algorithms
Informed Search

A Star (A^*) search

An agent is planning a path from start s to a goal g . Let's look at *node* n :

- UCS has the cost from s to n , which is the path cost $g(n)$, and it is the minimum in frontier
- It is backward cost from n 's perspective
- Greedy best-first considers only proximity to g , which is heuristic cost $h(n)$, it is minimum in frontier
- It is forward cost from n 's perspective



Figure 8: Consider both forward and backward cost while planning (image from internet)

A Star (A^*) search

Recap the pros and cons of UCS and Greedy search:

- UCS searches all directions \rightarrow no goal information
- It is optimal
- Greedy best-first is not optimal. Worst-case: badly-guided DFS
- Could takes the agent to goal straightly



Figure 9: Can agents compromise?

What about a new evaluation function $f(n)$?

$$f(n) = g(n) + h(n) \quad (2)$$

A Star (A^*) Search




Figure 10: State graph with spatial information

- A state graph with a start (S) and a goal (G), and other nodes
- Arrows indicate reach-ability, one direction
- Grids indicate spatial relationships \leftarrow Hops, Manhattan distance
- Adjacent grids may not reach each other directly, e.g., (S) and (b)

A Star (A^*) Search




Figure 11: Costs at node a

- Backward: path cost $g(a)$
- Forward: heuristic (hops) $h(n)$ from a to G
 1. Can be anything, i.e., number of blocks.
 2. Not necessarily real path cost from a to G
- $f(n) = g(n) + h(n)$ is an estimate of the full 'path cost'

A Star (A^*) Search - example




Figure 12: State graph with cost and heuristic

- Edge cost is the real distance
- Heuristic is hops. **Adjacency does not mean reach-ability**
- Can you calculate the hops and Manhattan distance, and what's the difference?

A Star (A^*) Search - example




Figure 13: Search tree with path cost $g(n)$ and heuristic $h(n)$

Next step?

1. UCS goes to ②, why?
2. Greedy search goes to ③, why?
3. A^* goes to ④, why?

Is A^* Optimal?




Figure 14: A sub-optimal path is picked by A^*

- UCS CAN pick the optimal route, **how?**
- Greedy search will not
- **Can we conclude the Greedy search contribution caused the issue?**

Admissible Heuristic

A heuristic $h(\cdot)$ is **admissible** (optimistic) if for a node n

$$0 \leq h(n) \leq h^*(n), \quad (3)$$

where $h^*(\cdot)$ is the true cost of a node to a nearest goal.





Figure 15: Admissible heuristic

- Main idea: estimated heuristic costs \leq actual costs
 - **Admissibility:** heuristic cost \leq actual cost to goal
 1. $h(A) \leq$ actual cost from A to G
 - **Consistency:** heuristic edge cost \leq actual cost for each edge
 1. $h(A) - h(C) \leq \text{cost}(A \text{ to } C)$
 2. **Triangle inequality:** $h(A) \leq \text{cost}(A \text{ to } C) + h(C)$
- Consequences of consistency:
 1. The f value along a path never decreases $h(A) \leq \text{cost}(A \text{ to } C) + h(C)$
 2. A^* graph search is optimal



Question: Is the condition $1 \leq h(.) \leq 2$ consistent?

Algorithm 1: Pseudocode of A^* Search

Input: Initial state s , goal state g , evaluation function $f = g + h$

Output: A *node* helps to retrieve a solution (path) \mathcal{P} , or failure

```
1: node  $\leftarrow$  with  $s$  as state
2:  $\mathcal{O} \leftarrow$  node, an priority queue           % node with the least cost.
3:  $\mathcal{C} \leftarrow \emptyset$ 
4: while  $\mathcal{O} \neq \emptyset$  do
5:   parent  $\leftarrow$  the first node in  $\mathcal{O}$        % 'first' due to the priority queue.
6:   if parent.state ==  $g$  then
7:     return parent
8:   end if
9:   del parent from  $\mathcal{O}$ 
10:   $\mathcal{C} \leftarrow$  parent
11:  for child in successor (of the current parent) do
12:     $v \leftarrow$  current child
```

```
13:    if  $v$  is not in  $\mathcal{C}$  and  $child$  is not in  $\mathcal{O}$  then
14:        add  $child$  to  $\mathcal{O}$       % found a new node.
15:    else if  $child$  is in  $\mathcal{O}$  then
16:        if current pathcost of  $child$  < previous pathcost of  $child$  then
17:            add current  $child$  to  $\mathcal{O}$ 
18:        end if
19:    end if
20: end for
21: end while
22: return failure
```







Figure 16: UCS

Figure 17: A^*

Graph search:

- A^* optimal if heuristic is consistent
- UCS optimal ($h = 0$ is consistent)

Consistency implies admissibility!

In general, most natural admissible heuristics tend to be consistent.